

CONTEXT AWARE SYSTEMS
A.Y. 2024/2025

Design and Implementation of a Geo-Aware Content Delivery Platform

Candidate:
Simone Tassi

0001186842

Introduction

- Modern mobile applications increasingly interact with the physical world, requiring systems that can **understand and react to user location in real time**
- Traditional cloud-based solutions often struggle to provide **low latency and scalability** in location-based, content-heavy scenarios
- Continuous location tracking raises **important privacy concerns**, requiring built-in mechanisms to protect user data
- This project presents a **geo-aware, edge based platform** that balances performance, scalability, and privacy

Introduction - Geofencing

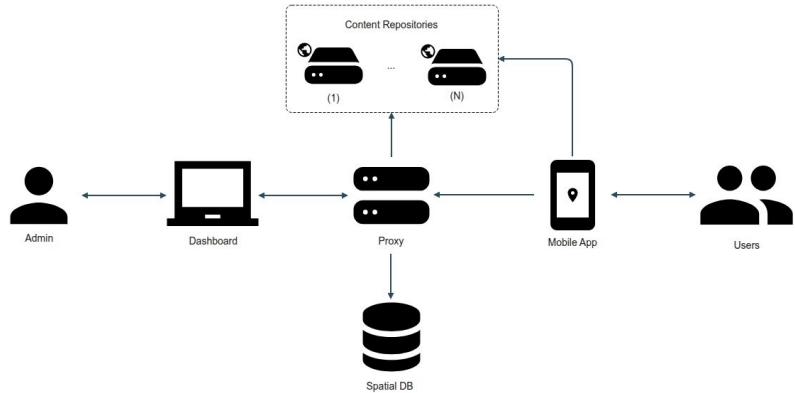
Geofencing is a technique that defines virtual geographic boundaries, allowing a system to detect when a user enters, exits, or moves within a specific area.

It is widely used in applications such as **location-based notifications, targeted content delivery, asset tracking, and smart city services**.

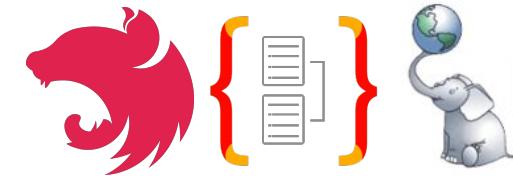
By continuously monitoring user location, systems can **trigger context-aware actions in real time**, improving user engagement and enabling personalized experiences.

Overview

- **Four main components:** a geo-aware proxy, distributed edge nodes, a mobile client, and a web-based dashboard.
- **Distributed Architecture:** a central proxy coordinating multiple edge nodes deployed across different locations to bring content closer to users.



Central Geo-Aware Proxy



Core Responsibilities

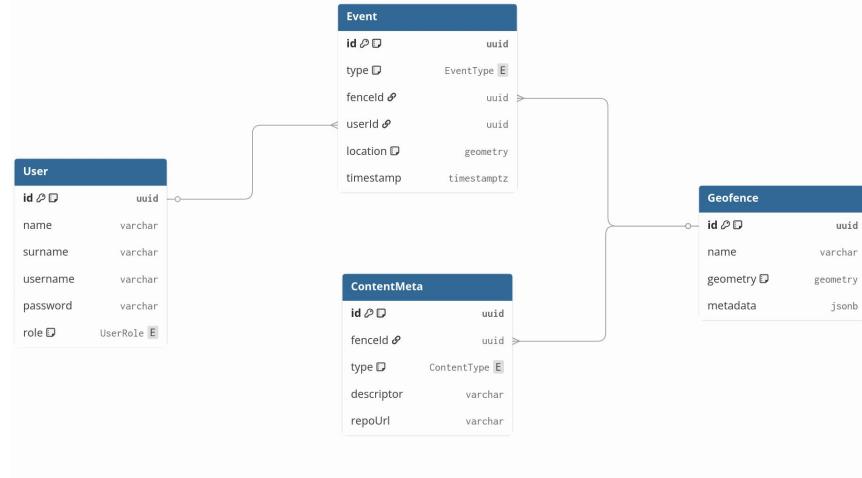
- Manages geofence definitions, content metadata, events and user identities
- Exposes REST APIs to web dashboard and mobile app
- On content request, issues a *302 Found* redirect toward the appropriate Edge Node, instead of directly delivering it to the client
- Integrates a real time WebSocket event bus to interact with the dashboard

Tech Stack

- NestJS, TypeORM, PostgreSQL (with PostGIS for spatial queries)

Data Model

- The **Geofence** is the central entity: it binds spatial boundaries to **ContentMeta** (one-to-many, cascade delete), with *repoUrl* pointing to the responsible edge node
- Events** capture **User** interactions in space and time, storing location as a native *PostGIS* geometry to power analytics and heatmaps downstream



Endpoints

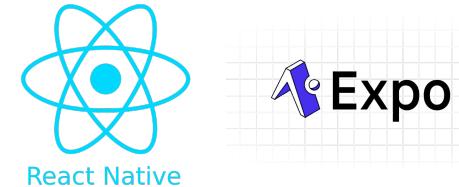
<i>Group</i>	<i>Paths</i>	<i>Operations</i>	<i>Methods</i>
Geofences	/geofences /geofences/{id} /geofences/active	CRUD + list active	GET, POST, PUT, DELETE
ContentMeta	/content-meta /content-meta/by-coords /content-meta/{id} /content-meta/content/{id}	CRUD + by coords + download	GET, POST, PATCH, DELETE
Users	/users /users/login /users/{id}	CRUD + login	GET, POST, PATCH, DELETE
Events	/events	create + list + wipe	GET, POST, DELETE
Analytics	/analytics/heatmap /analytics/metrics /analytics/clustering	heatmap + metrics + clustering	GET
Privacy Analysis	/privacy-analysis /privacy-analysis/simulate/{id} /privacy-analysis/export	simulate + export + wipe	GET, POST, DELETE

Edge Content Repository



- Lightweight **Fastify** nodes deployed close to physical landmarks
- Serving high-bandwidth assets (images, video, PDFs, ...) via **Node.js streams** → no full load into memory
- **SQLite** for local metadata persistence, keeping each node self-contained
- A **TTL mechanism** governs content freshness: expired assets return *410 Gone*, allowing the central proxy to filter out stale geofences automatically

Mobile App



Core Responsibilities

- Continuously monitor the user's position in the **background**
- Detect geofence entry/exit events
- Deliver event notifications
- Download content from the nearest edge node

Tech Stack

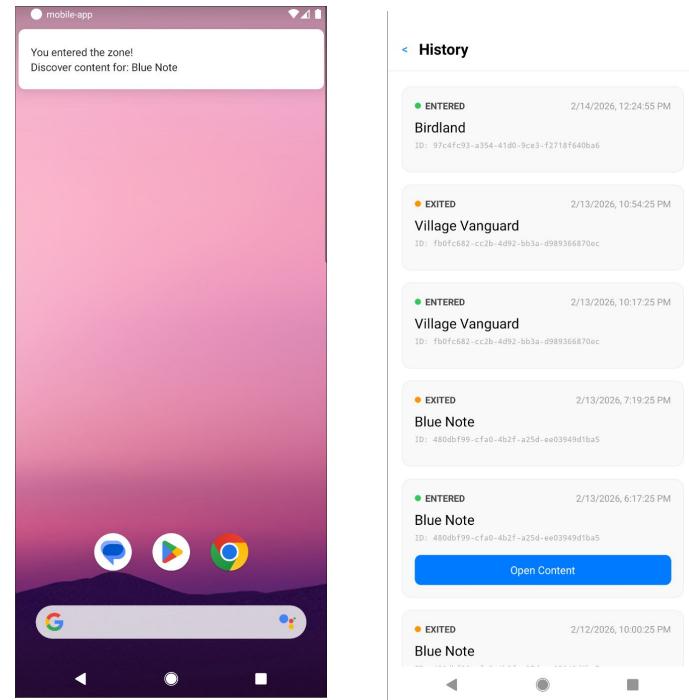
- **React Native + Expo** for cross-platform native capabilities
- **Zustand + AsyncStorage** for persistent state

Event Detection

Mobile OSes only support **circular region monitoring**.

Two tier hybrid strategy:

- **Tier 1 — Coarse:** low-power native circular monitoring as a proximity filter
- **Tier 2 — Precision:** high-accuracy GPS polling with **polygon containment checks locally** via *@turf/turf*, with no server involvement



Prefetching & Caching

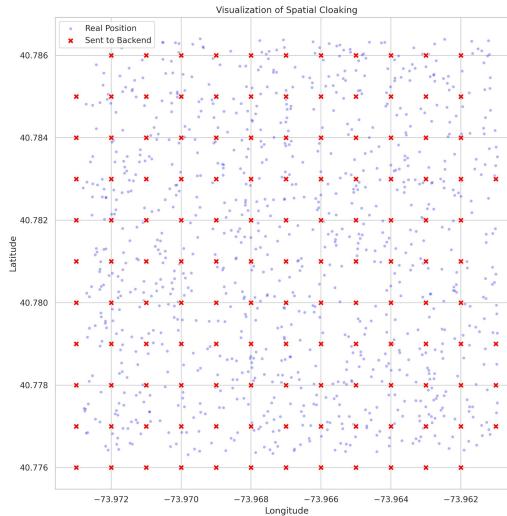
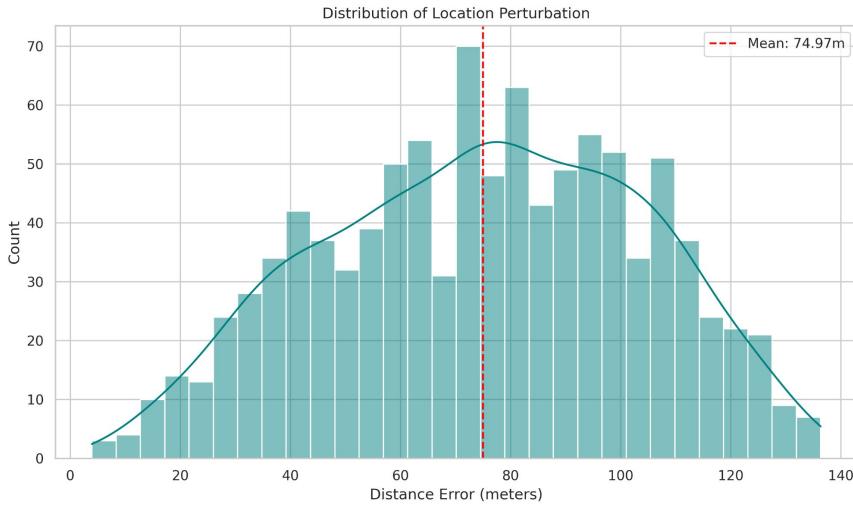
As soon as the user enters the **circular buffer** (*Tier 1*), content metadata and assets are downloaded in the background.

By the time the polygon is crossed and the notification fires, content is already available.

Assets are cached locally via **AsyncStorage**, enabling full offline viewing and resilience to connectivity drops.

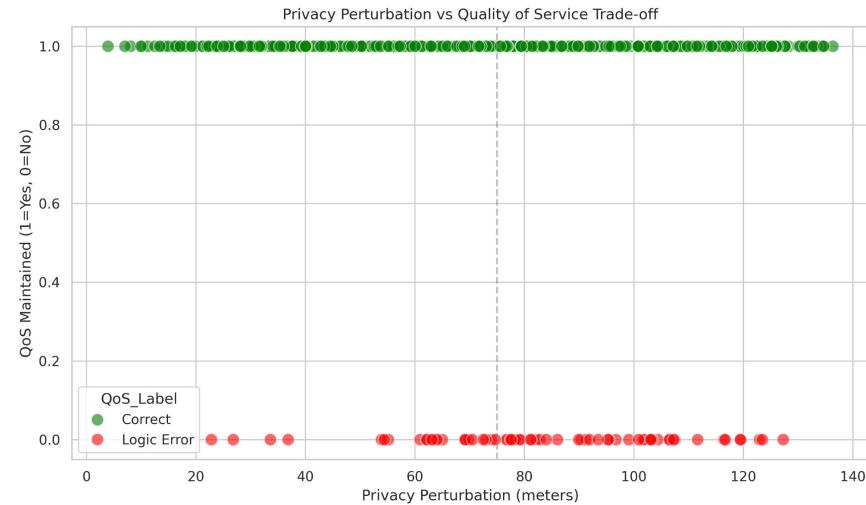
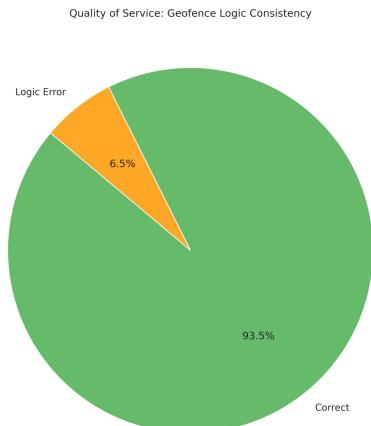
Location Perturbation

Users can toggle **Location Cloaking** from the home screen. When active, coordinates are snapped to a discrete grid ($\text{GRID_SIZE} = 0.001^\circ, \approx 110\text{m}$ at the equator) before being sent to the backend

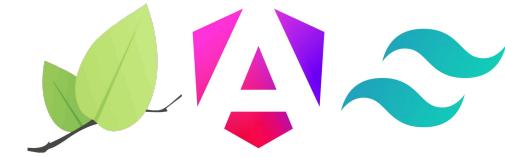


Location Perturbation

- The system maintains **high classification accuracy**, despite the spatial error introduced
- Errors are **not uniformly distributed**



Web Based Dashboard



Core Responsibilities

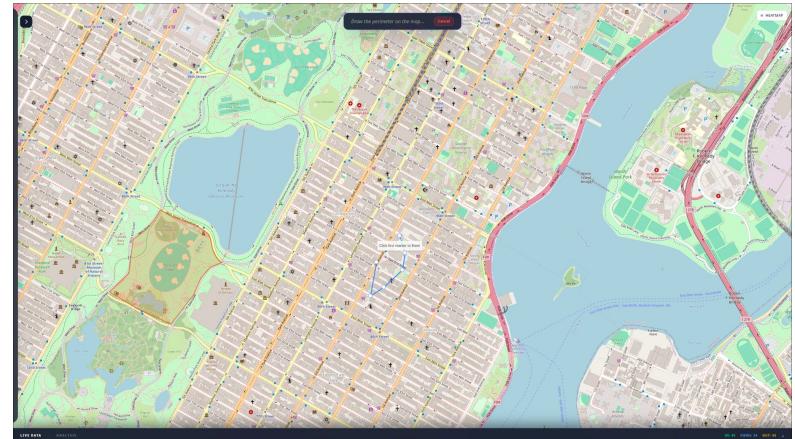
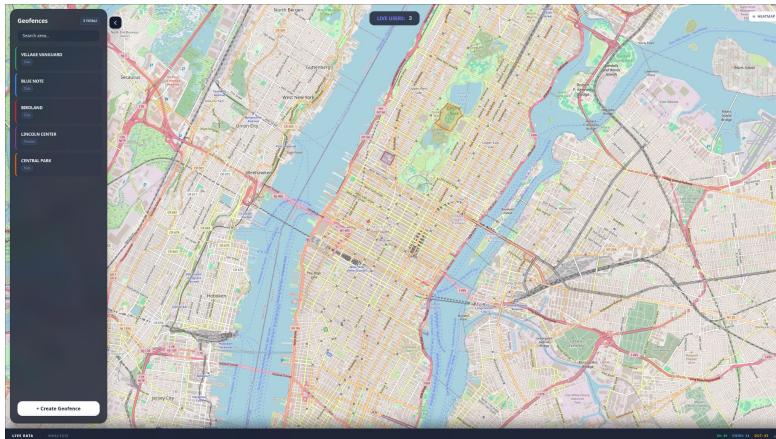
- Create and edit polygonal geofences via an interactive map editor
- Visualize live user activity and geofence events
- Show heatmaps of user engagement and clustering analytics

Tech Stack

- **Angular + Tailwind CSS** for a responsive UI
- **Leaflet + Geoman**: for polygon manipulation and geographical data visualization

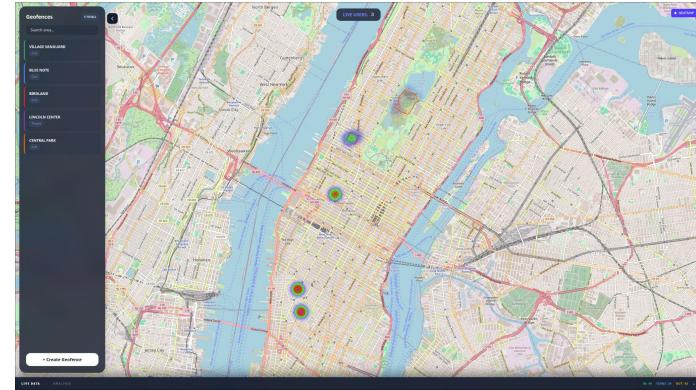
Geofence management

The ADMIN user is able to create, edit, visualize and filter geofences on an interactive map, enriching them with metadata such as *color* and a labelling system.



Live Analytics

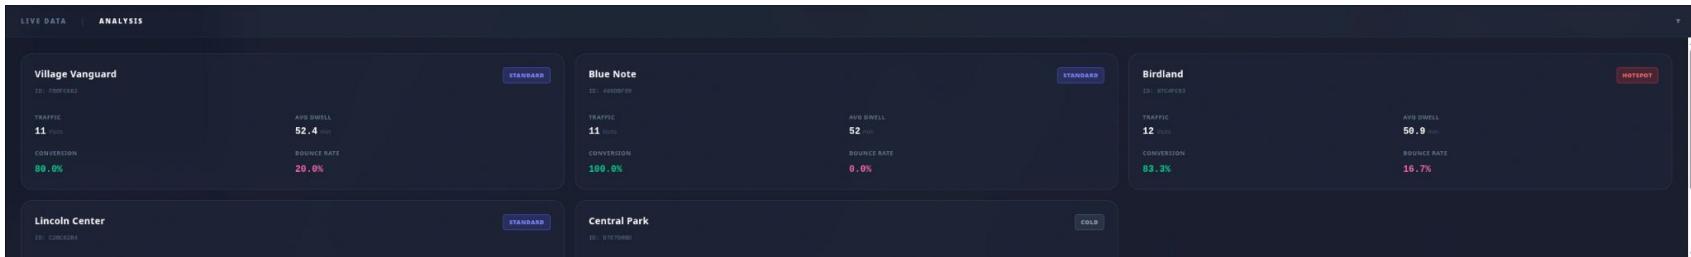
- **Real-Time Monitoring:** Visualize live user events through a WebSocket driven tray
- **Spatial Heatmaps:** Aggregate historical event data into a dynamic heatmap layer using *leaflet.heat*



Clustering - K-Means

- An asynchronous Python task uses *scikit-learn* to group geofences into three engagement levels: **Hotspot**, **Standard**, and **Cold**
- Clusters are automatically classified by ranking their centroids based on total entry volume

Feature	Description
<i>Entries</i>	Count of geofence entry events
<i>Conversion</i>	views / entries
<i>Dwell Time</i>	Duration spent by users within the geofence



Deployment

- The system is managed as a set of containerized services within a **Kubernetes cluster**, ensuring high availability and automated scaling
- A modular **Docker Compose** configuration is used for local development, providing a consistent environment across the entire development lifecycle
- Database persistence is handled via a **PostgreSQL + PostGIS StatefulSet**, utilizing **Persistent Volume Claims (PVC)** to ensure geographic data survives pod restarts

```
---  
apiVersion: apps/v1  
kind: StatefulSet  
metadata:  
  name: postgis  
spec:  
  serviceName: "postgis"  
  replicas: 1  
  selector:  
  template:  
    volumeClaimTemplates:  
    - metadata:  
        name: data  
      spec:  
        accessModes: [ "ReadWriteOnce" ]  
        resources:  
          requests:  
            storage: 16i  
  ---
```

Ingress

- An **NGINX Ingress Controller** acts as the single entry point, managing path-based routing for the REST API, frontend dashboard, and edge node endpoints
- Dedicated ingress configurations enable **Socket.io** tunnels for real time dashboard updates
- *Ingress* path rewriting rules allow for redirection of traffic to specific content repositories without exposing internal cluster networking

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: api-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /$2
    nginx.ingress.kubernetes.io/proxy-body-size: "50m"
    nginx.ingress.kubernetes.io/use-regex: "true"
spec:
  ingressClassName: nginx
  rules:
  - http:
      paths:
      - path: /repos/village-vanguard(/|$(.)*)
        pathType: ImplementationSpecific
        backend:
          service:
            name: service-village-vanguard
            port:
              number: 80
```

Node Affinity - Taint/Tolerations

- Use of ***nodeAffinity*** rules pins edge services to Kubernetes nodes physically located near the corresponding geofence
- Implementation of **Taints and Tolerations** ensures that dedicated edge hardware is reserved for localized content streaming, preventing interference from backend tasks

```
spec:  
  serviceName: "service-birdland"  
  replicas: 1  
  selector:  
    matchLabels:  
      app: content-repo  
      venue_name: birdland  
  template:  
    metadata:  
      labels:  
        app: content-repo  
        venue_name: birdland  
    spec:  
      affinity:  
        nodeAffinity:  
          requiredDuringSchedulingIgnoredDuringExecution:  
            nodeSelectorTerms:  
              - matchExpressions:  
                  - key: area  
                    operator: In  
                    values:  
                      - birdland  
          tolerations:  
            - key: "dedicated"  
              operator: "Equal"  
              value: "edge-node"  
              effect: "NoSchedule"
```

Conclusions

- Designed and deployed a full-stack platform that bridges the gap between centralized management and localized content delivery
- Implemented an automated analytics pipeline using *K-Means clustering* and real time *heatmaps* to transform raw location events into engagement metrics
- Integrated *caching* and *prefetching* strategies to ensure availability
- Proved that *location cloaking* mechanisms can provide a significant increase in user privacy with manageable impacts on *Quality of Service*
- Leveraged *Kubernetes orchestration* to ensure the system is resilient and highly available

THANK YOU